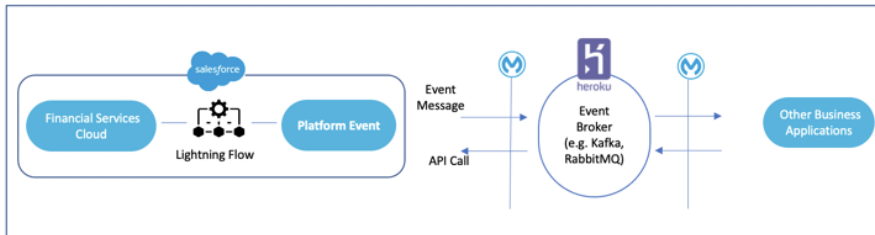


# Integrating Salesforce (Platform Events/Lightning Flow) with Kafka and MySQL: Hands-on Workshop

## Demo Scenario



- Account Address Change in Salesforce Financial Services Cloud triggers a Lightning Flow which generates a Platform Event
- A Mule Listener listens to the Platform Event and publishes it to Kafka on Heroku (after optionally transforming the Platform Event payload)
- Another Mule Listener consumes message from Kafka and routes it to a database connector which inserts the new address in MySQL database on Heroku.

## Steps

### PREREQUISITE

We assume that you have completed the FinsCabin Event-Driven Architecture workshop (Salesforce Event-Driven Architecture Hands-on Workshop: Student Workbook) and now have the demo environment and also now you are familiar with the various concepts (e.g. Mule Flows, Connectors & Transformations, Salesforce Platform Events, Lightning Flow & Heroku).

This demo assumes also assumes a basic working knowledge of Heroku CLI.

### Configure a new Mule Flow



Download this file and import it in Cloud Flow

Designer: <https://github.com/lightningexperience/finscabin/blob/main/resources/kafka/sf-events-kafka-db-1.0.0-mule-application-template.jar>

Please note that the following steps (in this sub-section - “Configure a Mule Flow”) are for your reference only. You should ideally download the above jar file as it contains Kafka-specific Mule DataWeave transformations to make it work with your existing Salesforce Account Address Change flow and the Heroku FinsCabin Heroku app.

Here is a brief explanation of how I created a new Mule flow for Kafka from the existing Event-Driven Architecture Workshop flow which I had originally used for RabbitMQ (instead of Kafka).

Clone your existing Mule demo flow which connects Salesforce to Database. For this, you can save your existing demo as a template and then use this template to create a new Mule flow.



Now in the new Mule flow, that you have created from your template - delete AMQP connectors. You can optionally rename the sub-flows in the Cloud Flow designer(e.g instead of Salesforce to AMQP you can rename it as Salesforce to Kafka

In the Cloud Flow Designer insert the Kafka Publish Connector where you had AMQP Publish Connector.




In the AMQP to Database subflow - remove AQMP and insert Kafka Listener



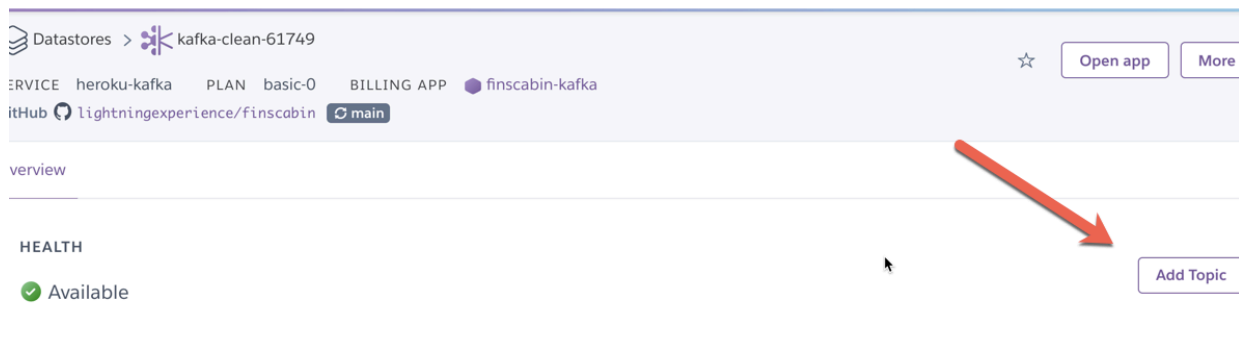
We will now need to connect these Mule Connectors to Kafka - let's first install Kafka on Heroku.

## INSTALL & CONFIGURE KAFKA ON HEROKU

Go to your Heroku instance and provision Kafka in your existing FinsCabin app. **Be aware that Kafka does not have a free tier on Heroku** - please check with your manager before installing.



Now create a topic in Heroku (you can do it through the Heroku GUI). Let's call it eda01. It will also ask you for number of partitions (Just select 1).



The topic needs nearly 30 minutes to get created and become visible in Heroku CLI (You won't see topics listed in the dashboard).

Install Heroku CLI - You would need it in this demo: <https://devcenter.heroku.com/categories/command-line>

Login to Heroku via CLI

```
heroku login
```

Install Kafka Plugin for Heroku CLI

```
heroku plugins:install heroku-kafka
```

Create a Consumer Group in your Kafka instance.

```
heroku kafka:consumer-groups:create consumer01 -a finscabin-kafka
```

Now you can view your Kafka instance from CLI

```
heroku kafka:info -a finscabin-kafka
```

## CONNECT TO HEROKU KAFKA FROM MULE CONNECTOR

We are making this assumption that your environment is in Heroku multi-tenant Common Runtime (not Private Spaces). In Private Spaces you have some more options for connecting including direct access to Zookeeper. Let's keep things simple and assume you have the simplest Heroku configuration (default in Personal/demo/trial accounts).

Heroku provides the certificates to connect to Kafka in PEM format. The MuleSoft Kafka Connector requires the certificates to be in JKS format. Therefore, we need to convert the PEM certificates and keys to a PKCS12 file first before we convert them to JKS.

Go to your Heroku app -> Settings and view the config variables.

You will see certificates as in the following screenshot.



Copy and paste them in individual files using a text editor choosing your respective file names (e.g. cert.pem) as I have done below.

- KAFKA\_CLIENT\_CERT -> cert.pem
- KAFKA\_CLIENT\_CERT\_KEY -> key.pem
- KAFKA\_TRUSTED\_CERT -> trusted\_cert.pem

Save all the three files in one directory and then navigate to that directory from command line.

Run the following command to generate a **pkcs12** file from the **cert.pem** and **key.pem** file:

```
openssl pkcs12 -export -out cert.pkcs12 -in cert.pem -inkey key.pem
```

The command will ask you for an export password. Don't forget to write that down. You'll need it for the next step as well as for the configuration of the Kafka Connector.

Once that's done, run the following tool to convert the **pkcs12** file to a JKS.

```
keytool -importkeystore -srckeystore cert.pkcs12 -srcstoretype pkcs12 -destkeystore keystore.jks -deststoretype jks -deststorepass mypassword
```

Make a note of all the passwords you use in the above commands - you will need it in the Mule Kafka Connector configuration.

## Create Truststore

1 - Truststores require the certificate to be in binary format. The current format exported from Heroku Kafka is PEM format. Run the following command to convert the PEM certificate to binary format (DER). This creates the binary file with the name cert.der

```
openssl x509 -in trusted_cert.pem -out cert.der -outform der
```

2 - Import the binary certificate into a new jks truststore called truststore.jks by running the following. When prompted, type yes.

```
keytool -importcert -file cert.der -keystore truststore.jks -storepass <your-key>
```

### Output

Only the last 2 lines of the output are shown for brevity.

```
Trust this certificate? [no]: yes
Certificate was added to keystore
```

The keystore.jks and the truststore.jks files are what we need for the next section. In my connector configuration I am using heroku01.jks as the keystore file name and truststore01.jks as the trustore file name.

## CONFIGURE MULE CONNECTORS TO CONNECT TO KAFKA

In the Kafka installed in the Common Runtime (multi-tenant) Heroku you need to prefix the topic and consumer name with the Kafka Prefix environment variable which you will find in your Heroku app's setting.

KAFKA_PREFIX	tennessee-75801.

When putting in the name of topic in your Mule Kafka connector, you will include the suffix - e.g. tennessee-75801.eda01 instead of simply eda01. This also goes for the consumer group (e.g. tennessee-75801.consumer01 instead of simply consumer01).

Open the Kafka Publish Connector and put in the value as shown below.

Apache Kafka Connector  
Publish

Configuration Input Output Not

ing the Apache Kafka Producer configuration configuration. [Edit](#) [Search configurations](#)

General Topic ⓘ

Advanced tennessee-75801.eda01 ⓘ

Partition ⓘ

CUSTOM EXPRESSIONS

Custom Expression

PAYLOAD

name

Now edit to add a connection.

The screenshot shows the 'Salesforce to Kafka' configuration interface. At the top, there's a header with the Salesforce logo and 'Apache Kafka Connector Publish'. Below this are tabs for 'Configuration', 'Input', 'Output', and 'Notes'. The 'Configuration' tab is active, showing a message: 'Using the Apache Kafka Producer configuration configuration.' with links for 'Edit' and 'Switch configurations'. A red arrow points to the 'Edit' link. Below the message, there are two sections: 'General' and 'Advanced'. The 'General' section has a 'Topic' field with the value 'tennessee-75801.eda01' and a 'Partition' field. The 'Advanced' section is currently empty. On the right side, there are sections for 'CUSTOM EXPRESSIONS' and 'PAYLOAD', each with a 'Custom Expression' field and a 'name' field.

Copy the Heroku "KAFKA\_URL" from your apps' setting page and put it in the Bootstrap Server URL field as shown below. Fill in other details as below.

The screenshot shows the 'Configure Connection' dialog box. It has a title bar 'Configure Connection'. Inside, there's a 'Connection' section with a 'Connection Name (required)' field containing 'finscabinkafka'. Below this is a checkbox labeled 'Share this connection with my business group' which is checked. Underneath is a 'Connection Type (required)' dropdown menu showing 'Producer Plaintext Connection'. Below the 'Connection' section is a 'General' section with a 'Bootstrap Server URL (required)' field. At the bottom of the dialog are three buttons: 'Test', 'Cancel', and 'Save'.

General

Bootstrap Server URLs (required) ⓘ

kafka+ssl://ec2-100-25-107-37.compute-1.amazonaws.com:9096,kafka+ssl://ec2-3-220-121-33.compute-1.amazonaws.com:909

General

TLS Configuration

Enabled Protocols ⓘ

TLSv1, TLSv1.1, TLSv1.2

Test Cancel Save

## Configure Connection

Enabled Cipher Suites ⓘ

Trust Store

Path

truststore03.jks x v Upload

Password ⓘ

..... Show

Type ⓘ

Test Cancel Save

Password ⓘ

..... Show

Type ⓘ

JKS

Algorithm ⓘ

I

## Configure Connection

**Your file name may be keystone.jks**

Key Store

Path

heroku0103.jks x v Upload

Type ⓘ

JKS

Alias ⓘ

## Configure Connection

**To keep things simple I used the same password when generating the files**

Key Password ⓘ

..... Show

Password ⓘ

..... Show

Algorithm ⓘ

Delegation Check

Test Cancel Save

For the endpoint identification algorithm I using space (i.e. double quotes with nothing between them as in the following screenshot).



A screenshot of a configuration interface. At the top, there is a dropdown menu with the text "None" and a downward arrow. Below this, a section titled "Endpoint Identification Algorithm ⓘ" is highlighted with a red hand-drawn oval. Under this title, there is a text input field containing the value "#[\""]".

Now test (you should get success). Save and again Save.

Now let's configure the other sub-flow - the Kafka Consumer Connector's connection information to Kafka.

A screenshot of a configuration window titled "Kafka to Database". The window has tabs for "Configuration", "Output", and "Notes". The "Configuration" tab is active. Below the tabs, it says "Using the Apache Kafka Consumer configuration configuration:" followed by links for "Edit" and "Switch configurations". A red arrow points from the top of the window down to the "Edit" link. On the left side, there is a sidebar with "General", "Redelivery", and "Advanced" sections. The "General" section is selected, showing "Poll Timeout ⓘ" and "Poll Timeout Time Unit ⓘ" with a dropdown menu set to "Select an option".

## Configure Connection

A screenshot of a "Configure Connection" dialog box. It has a "Connection" section with a "Connection Name (required) ⓘ" field containing the text "finscabin-kafka01". Below this is a checkbox labeled "Share this connection with my business group" which is checked. Underneath is a "Connection Type (required)" dropdown menu with "Consumer Plaintext Connection" selected. A red arrow points from the top of the dialog down to this dropdown. At the bottom of the dialog, there is a "General" section with a partially visible "Bootstrap Servers (required) ⓘ" field. At the very bottom, there are three buttons: "Test", "Cancel", and "Save".

General

Bootstrap Server URLs (required) ⓘ

kafka+ssl://ec2-100-25-107-37.compute-1.amazonaws.com:9096,kafka+ssl://ec2-3-220-121-33.compute-1.amazonaws.com:9096

Group ID ⓘ

tennessee-75801.consumer01

General

TLS Configuration

## Configure Connection

Path

truststore04.jks

Upload

Password ⓘ

.....

Show

Type ⓘ

JKS

Algorithm ⓘ

## Configure Connection

Algorithm ⓘ

☒ Insecure ⓘ

**Your file name may be keystore.jks**

Key Store

Path

heroku0104.jks

Upload

### Configure Connection

heroku0104.jks

x v

Upload

Type ⓘ

JKS

Alias ⓘ

Key Password ⓘ

.....

Show

Endpoint Identification Algorithm ⓘ

#[""]

Now as you continue filling in the connection info in the Kafka Listener Mule connector - click on Advanced and scroll down to the bottom. Fill in the topic name as shown below.

Seconds v

Topics

Topic Subscription Patterns ⓘ

tennessee-75801.eda01

Add

Assignments ⓘ

Add

Now test (you should get success). Save and again Save.

The rest of the workshop would work as explained in the Event-Driven Architecture Student Guide: Salesforce Event-Driven Architecture Hands-on Workshop: Student Workbook

## Resources

- Salesforce Event-Driven Architecture Hands-on Workshop: Student Workbook
- How to connect to Heroku Kafka from Mule - In my example above I have used the compact keytool command and did not need Jetty as this blog advises: <https://blogs.mulesoft.com/dev-guides/api-connectors-templates/how-to-connect-to-apache-kafka-on-heroku/>